

Table of Contents

Summary	1
Target Audience	2
Introduction	2
Overview: Key Concepts and Terminology	2
Putting it Together: Creating Modern AI and Machine Learning	4
Running AI Projects	5
Using AI	5
AI Project Decisions	5
AI Project Workflow	5
Data	5
The Importance of Data	5
Collecting Data	5
Managing Data	5
Validation, Cross Validation and Data Leakage.	5
K-Fold Cross Validation	5
Train, Validate and Test	6
Building AI and Machine Learning Models	7
Decisions	7
Decision Making Flowchart	9
Neural Networks	9
Worked Examples	12
Pitfalls and Common Problems	12
Resources	12
Appendix	12

Summary

Artificial Intelligence (AI) and Machine Learning (ML) technologies are key to many modern engineering projects due to their ability to solve many problems that are difficult or impossible with other methods. While most engineers will find themselves enjoying a significant overlap between these techniques and their existing skill set, they are also liable to find that AI and Machine Learning is its own field with its own unique demands and (often hidden) pitfalls. While there are many resources available for self teaching, it is generally assumed the practitioner is either an absolute beginner to engineering, or already a seasoned expert in AI and ML. In this document, we provide a practical guide to AI and Machine learning for electronic systems engineers who already have a strong base of knowledge in electronic systems but no specialised expertise. This guide will be practice focused, with the goal of helping engineers to make good decisions and avoid problems. The guide will cover, among many others areas:

- Technical Language and Core Concepts
- Data, Collection and Common Problems
- Deploying AI and Machine Learning, with Worked Examples

- Existing resources, where to find them and how best to use them
- Common Pitfalls, how to avoid and how to solve

Target Audience

The target audience for this guide is electronic systems engineers with little to no specific expertise in AI and Machine Learning techniques. It will be assumed, because of this background, that readers will have a base level of competence in programming and mathematics, though the guide will err on the side of caution in respect of assumed knowledge. We may, for example, assume our reader has knowledge of concepts in basic calculus, but are unlikely to assume that they are able to remember any specific formula. We are also assuming that the primary interest in practically deploying these techniques rather than understanding the theory and context of their development. For readers interested in a more theoretical treatment, we list several texts in the resources section appropriate to a range of different levels of background knowledge.

Introduction

The terms Artificial Intelligence and Machine Learning both lend themselves to the idea that we are creating a program that can, in some respect, think and learn in the same way a person can. While it may be possible to create systems that are capable of this, the current technologies that are deployed in practice are much more restricted in scope. Current, practical Artificial Intelligence and Machine Learning technology generally refers to programs that are of a statistical, or statistically adjacent nature, and that use data to fit some specified mathematical model. While these techniques are proving extremely valuable in solving otherwise difficult or intractable problems, the fundamentals of the techniques are neither exceptionally novel nor complicated. In this guide, we aim to show just how simple it can be to deploy these techniques in practice, especially for an audience that already has knowledge of electronic systems engineering.

It is important to acknowledge that, despite the above, no specific definition of Artificial Intelligence and Machine Learning that is universally agreed on exists. Furthermore, lacking any significant advancement in our understanding of what cognition, intelligence, and learning actually mean, this state of affairs is likely to continue. Given the practical focus of this guide, it would not be profitable for us to enter into this debate at any length. We will therefore, universally use the term “AI and Machine Learning” to refer to algorithms that are the subject of this guide, and that learn to **parameterize** a mathematical model from data.

This guide is structured into 5 sections. In the first section, we discuss the key concepts and terminology that are used in AI and Machine Learning and briefly discuss how they come together to motivate us to create the field in its current form. With this important background covered, we then move on to discussing the practicalities of AI and Machine Learning projects at a high level. We discuss where it's appropriate to deploy these technologies, a high level workflow and decision making process, and what management of these projects looks like. Having covered this, we then begin to drill down into the details of these projects. We open this discussion with a review of the key role data, and how it can be managed. This is followed with an extensive discussion on the implementation and decision making process of deployment of these algorithms with worked examples. We close with a list of pitfalls, common problems and solutions, and a link to useful further resources.

Overview: Key Concepts and Terminology

In this section, we introduce and overview some of the key concepts and terminology that are important in understanding and discussing AI and machine learning technology. This section is not intended to serve as an exhaustive discussion of the topics in question, but as an introduction to the key points and vocabulary that will frame the rest of the discussion on this topic. Key terms and concepts will be **highlighted** for the reader, with links to a more detailed description of the term in the appendix.

To briefly review what we discussed in the introduction: the type of intelligence and learning in algorithms that we are interested in in this manual are algorithms that learn mathematical models of the world from

some data, in order to make predictions about other unseen or future data. One important idea that we need to consider first is **structured data** and **unstructured data**.

Breakaway: Structured vs Unstructured Data: AI and Machine Learning models are no different from any other computer program in that they require their input data to follow a consistent format. Unfortunately, data collected in the real world rarely follows the type of structure and organization that is needed for ingestion by an AI or Machine learning algorithm, and in many cases the work done to transform data into an appropriate structured format is some of the most important work done in any AI and Machine Learning pipeline. We make this distinction between data that has been put into a useful structured format as **structured data**, and data that exists in a raw, unprocessed format as **unstructured data**.

When dealing with data in the real world, we will often split it up into categories or types. One such distinction often made that is especially important in the context of AI and Machine Learning is the split of data into **continuous** data and **discrete** data. Continuous data can take on any number of infinite values across a given range, for example, a measure of rainfall per hour. Discrete data on the other hand is any type of data that falls into a fixed number of categories. These categories can be both **ordinal** data in which there is a natural ordering between the categories (shoe size, for example), and **nominal** data, where the categories are distinct (eye color, for example). While this distinction is important for many parts of AI and Machine Learning, the distinction between whether an AI and Machine Learning algorithm is trying to predict continuous and discrete data is so important that it has its own nomenclature of **regression** and **classification** algorithms respectively.

Breakaway: Regression vs Classification Algorithms: * The distinction between **regression** (continuous output data) and **classification** (discrete output data) is particularly important in AI and Machine Learning algorithms, because the type of data that the algorithm outputs has a significant effect on how it must function. Notably, some algorithms (e.g. Support Vector Machines) are only designed to function in one of these modalities, and require significant adaptations to perform (likely very poorly) in the other.

While we have been discussing some of the concepts and terminology around data to this point, we have used the terms “learn”, “learning” and “learning from data” to describe what our algorithms do without really making it explicit what we actually mean by this. One of the reasons that we’ve avoided doing this is that “learning” in the context we’re discussing it is conveniently, without further qualifiers, a term that covers several different ideas. These differences stem from the way that we use data in order to “learn”. The most prominent of two of these ideas are **supervised learning** and **unsupervised learning**, which are concerned whether we learn from data that list the correct output the algorithms should produce for some given input data (**labeled data**), or simply the input data themselves (**unlabeled data**).

Breakaway: Supervised vs Unsupervised vs Reinforcement vs Other Learning: We use the nomenclature of **Supervised** vs **Unsupervised** (vs others) to describe the way in which our algorithms are learning. In Supervised learning, we learn from matched input data/output data pairs, data for which we already have the correct output the algorithms should predict for a set of given inputs (“learning by example”). We call this data **labeled data**, because our set of input data is labeled with the corresponding correct solutions. For example, we might be interested in predicting the future prices of the stock market from economic indications, by looking at how these economic indicators have predicted its historical past prices. In Unsupervised learning, we only have access to the input data without any corresponding output solution attached. We call this data **unlabeled data**, and our unsupervised learning algorithms are generally interested in predicting some quality of this data (“pattern learning”). For example, we might be detecting unusual anomalies of electrical usage in the grid.

While it is generally preferable to use supervised learning when we can because learning by example is easier, there are many situations in which unsupervised approaches are more appropriate. Even putting aside the fact that unlabeled data is easier to collect (since we don’t need to label it), for many problems supervised approaches are simply not practical. In our electrical grid example above, it would be infeasible to train a supervised model to do similar anomaly detection. By definition, anomalies are rare and unusual data points that fall outside of the usual observations in the data. Creating a labeled dataset of them would be both impractical, and any supervised algorithm that used it would be prescriptive - it would only catch anomalies similar to anomalies we’ve trained on, where an unsupervised approach instead catches ones that are dissimilar to everything we’ve seen so far.

There are also several other learning approaches that fit within the supervised/unsupervised dichotomy discussed so far. A common one is **Reinforcement Learning**. In Reinforcement Learning, the algorithm is not fed a set of data, but selects which piece of data it wants to learn from in future from the pieces of data it has had up until now. Another common paradigm is **semi-supervised learning**, in which an algorithm learns from some set data that is labeled, and some (usually larger) set of data that is unlabelled.

No matter which of these learning types we want to use, we need to be able to evaluate the performance of the AI and Machine Learning models we create. The way we approach this is no different to any other testing we would do - we compare the predictions that our model makes to some known ground truth data. An easy way to do this would be, once we have **trained** our algorithm on the data that we have to hand, to test how well it performs on this same data (evaluate the **training error**) as a ground truth. Unfortunately, this is a bad idea. AI and Machine Learning algorithms will fit fairly well to the data they've trained on (it's **training data**), independently of how well they work for other "unseen" data. Since the goal of our algorithm is to have it work well across all data points (including ones it wasn't trained on), how well it performs on the training data will be a misleading and overconfident measure of overall performance.

Instead, we try to estimate how well our model will perform on data we've not trained on by randomly reserving a small amount of our data in a testing set (our **testing data**). Sometimes, in addition to the **training** and **testing** sets we've described, we will make a further split of our data to also include a small set of **validation** data. We might do this if we need to validate the results of testing, for example, in more advanced applications in which we might use the testing data itself to make decisions about the learning process. It is generally best practice not just to break your data up into training and testing (and validation, if needed) sets once, but to repeat this process multiple times and aggregate the results. This process is called cross validation, and in most cases will be the more appropriate way to evaluate our AI and Machine learning model's performance.

Breakaway: Measures of Performance: * There are multiple ways for us to evaluate performance of any given model. Some common choices are ***Mean Squared Error** (continuous data) or **Cross Entropy** (discrete data). It's best to stick to standard measures unless you understand what you're doing, but there are usually multiple valid ways of measuring performance with their own consequences. The best measure of performance is the one that solves your problem best.

With a measure of performance, we have a way of comparing different models to select the best one. Practically though, there are too many different algorithms and approaches for us to run them all and directly compare them in this way. We need a way of selecting likely candidates a-priori, without directly testing them. Our goal in AI and Machine Learning is to make predictions about all of our data from a small subset of it. We want a model that accurately reflects the reality of the data we're training it on. An abstraction that can help us think about this is to think in terms of **model complexity**. Our models exist broadly on a spectrum of complexity from simple linear models with only a few parameters that fit a line to our data at one end, and billion-parameter neural networks at the other. It's probably clear to see that an overly simple model of our data will be bad. If we can't capture the complexity of what is happening in our data, we'll never be able to model it well. We call this **underfitting**. However, it's also the case that fitting a model that is too complex is problematic. Models that are too complicated will fit randomness in the specific data they are trained on, and will not generalise well to data outside of that. We call this problem **overfitting**.

Our goal should be to pick a model that is complicated enough to fit to the parts of the data we are interested in (the **signal**), without overfitting to the noise in our data too. We also want to take advantage of prior knowledge we have about our problem, for example, if we know our problem is linear, it would be sensible to pick a linear model. When in doubt, it's often more favourable to go for simpler models, for reasons we will discuss later.

Putting it Together: Creating Modern AI and Machine Learning

WIP

Running AI Projects

In this section, we discuss the problem of designing and managing an AI and Machine Learning project. Importantly, this is *not* a technical guide to solving these problems, but a guide to solving all the problems that precede and surround the technical parts of the problem.

Using AI

The first, and most important problem to solve in any AI and Machine Learning project is to be able to formulate a clear and concise answer to the question “why do I want to solve this problem with AI and Machine Learning”? AI and Machine Learning algorithms are far from universally appropriate solutions, and suffer from several fundamental difficulties that make them undesirable:

- They require collection and processing of data to feed them
- They are stochastic, dealing fundamentally in probabilities
- They are difficult to validate, and further yet, many algorithms are difficult even to interpret

The reason these algorithms have received so much attention *despite* these difficulties is that they make it plausible (or possible) to solve sets of problems that are otherwise difficult to get at. These challenges, and the motivation we gave for AI in our previous sections (understand a large data set by learning from a (relatively) small amount of data) speak to a litmus test for whether a problem is suitable to be solved with AI and Machine learning. A problem is a good candidate if:

- It is infeasible to solve the problem in a more direct or analytical way
- It is feasible to access a useful set of data points to indirectly learn a solution
- It is infeasible to access all (or almost all) of the data points we're interested in

AI Project Decisions

AI Project Workflow

Data

The Importance of Data

Collecting Data

Managing Data

Validation, Cross Validation and Data Leakage.

At some point we are going to need to assess our models' performance on unseen data, as this is the only way to get an unbiased estimate of its performance. There are two popular ways this is done.

K-Fold Cross Validation

1. **Randomly** split the dataset into K **equal** partitions.
2. Randomly initialise the model and train on all-but-one of the partitions.
3. Test model performance on the withheld partition.

4. Repeat steps 2 and 3, withholding a different partition each time, until all K partitions have been used as the test data.
5. Combine the K accuracy scores at the end, providing a mean and variance for the accuracy.
6. Make changes to model and repeat until performance is sufficient.

This way, all of the available data is used for both training and testing, while never testing a model on the data on which it was trained. It also serves to avoid any unintended bias that may arise by chance, when the data is randomly partitioned. However, it requires training the model from scratch K times to get a single output. Also, repeatedly running K-Fold Cross Validation and making adjustments, can introduce implicit bias.

This is well suited to training smaller models on small datasets, as the requirement to train many times rules out its use in 'big data' applications such as deep learning. If data is limited and your model can be trained many times (at an acceptable cost for you), K-Fold Cross Validation is recommended. Otherwise, read on to the next section.

Train, Validate and Test

1. **Randomly** split the dataset into 3 partitions. The proportions of these are up to you, but a sensible split usually looks something like 80%-10%-10% or 60%-20%-20%, (train-validate-test, respectively)*.
2. Repeat until performance is sufficient:
 - a. Train the model on the train set.
 - b. Evaluate on the validation set.
 - c. Make any changes to the model.
3. Evaluate your best performing model(s) on the test set.

This should be the end of the process, with these results being publication/report ready. One should avoid going back and making more changes after this, as that is how implicit bias from the testing set can creep into the model design (data leakage).

*The larger the datasets, the less variance we can expect between them. So if we expect to do lots of hyperparameter tuning iterations (step 2), we might opt for a 60%-20%-20% split. That way the information we gain from the validation step should be more dependable. If not, the 80%-10%-10% would be preferable simply as we are using more of the data for training, which should give the best model performance.

Points to remember:

- We optimise for performance on the training data, but we **assess** performance on 'unseen' (not training) data.
 - If it can perform well on this, we know the model isn't simply 'remembering' the correct answers from its training process but has learned meaningful relationships which apply outside of its training dataset.
 - We can think of this the same way we might think of assessing the learning ability of a student. That being, they 'train' on practice questions and are assessed on 'test' questions which require the same skills as the training questions, but the students have never seen them before. Assessing them on questions they have already seen (training data) clearly would not be a sensible way to evaluate their understanding of the topic.
- Unseen data can still influence the model, via influencing the **decisions we make** to improve the model. This is called **data leakage**.
 - This is a very subtle problem which can easily go undetected.

- To avoid this, make the train/validate/test as soon as possible, and leave the test dataset untouched until it's time to evaluate our model. This includes not visualising or exploring the test data, to prevent bias in our own thinking.
- This can again be likened to a student taking a test - if they know what questions are coming up and prepare for those specifically, their test score will end up higher than a true reflection of their ability.

Building AI and Machine Learning Models

Decisions

Data Pipeline

Data Pipeline is simply a term used to describe the moving and processing of data from one place to another. While this is a broad term, the case we are concerned with specifically is loading and transforming data from storage in order to be used by a neural network.

Train, test, validate: See earlier section for details. Building the data pipeline is where we must take this into account and avoid the common pitfalls listed in this section.

Reading the data: Whichever machine learning API we use, three core functions typically need to be implemented for a custom pipeline:

- Initialise.
- Get dataset length. In order to know when an epoch is complete, the class needs to know how much data there is in total. This function is as simple as it sounds.
- Return the i th datum. A function with a one to one mapping of integers i to data. This is important so that the data loader can shuffle the dataset by shuffling a list of integers from zero to dataset length.

Transformations. Transforming data is an important step in machine learning. Normalising data is usually beneficial in this case, but there may also be transformations specific to your task. A decision that should be made early on in this process is when to do these transformations. The options are as follows:

- Add transformations to the data loading step.

Pros:

- Simple, readable
- Easy to change
- Minimal code

Cons:

- Each transform must be applied many times (every time the data is loaded, once * per epoch per datum). This can be slow.

- Apply transforms to the entire dataset and save in storage.

Pros:

- Faster
- Only needs to be considered once

Cons:

- Harder to quickly experiment with different transformations
- Requires either overwriting original data, or doubling storage requirements

TechNES AI Best Practices Group:

- If computational limitations (sufficient RAM:dataset size ratio) allow, there is an optimal middle ground. This is to use the first method but load all the data into RAM before training.

Pros:

- All the pros of the above methods
- In addition to avoiding repeated applications of transforms, it also avoids potentially slow read times from storage, in favour of fast reading from RAM.

Cons: * High RAM requirements, which increase with dataset size. * Even if HPC allows this, it cannot be tested on local machines with less RAM, without modification.

Batch Size: The amount of data the model sees before optimising its parameters at each step. Larger batch sizes will give a better estimate of the true optimal parameter adjustment the model should make, but each step will take longer to compute. There is no universally optimal number for this. It is advisable to use powers of 2, to aid with parallelisation and GPU usage in case this is desired later. Consider starting with a batch size of 32 or 64 and increasing/decreasing if the training is unstable/convergence is too slow, respectively.

Epochs: This is the number of full passes through the dataset that the model will see before training finishes. It can either be a fixed number, or an upper bound, only reached if the model does not converge. The latter scheme is preferable for maximising performance, since early stopping (end training as soon as convergence is reached) is known to be a form of regularisation which can improve performance on unseen data. The former may be used for comparing training stability of multiple models.

Training Process

There are a few more decisions to make before we get to designing and training custom neural networks.

Loss Function: This is the quantity that the optimiser will try to minimise, thus it is critical to choose something that, when minimised, makes the model more useful. This function will be applied to the output of the network and the target. For example, the mean square error loss function calculates the mean of the squared error between the output and target, thus minimising this quantity forces the model to try to improve prediction accuracy. This is particularly useful for regression problems in which every element of the output is equally important. However, if we want a model which predicts a probability distribution, we may choose to use the categorical cross entropy or KL-Divergence for example. A detailed breakdown can be found here <https://machinelearningmastery.com/how-to-choose-loss-functions-when-training-deep-learning-neural-networks/>.

Optimiser: At every step in the optimisation process, the algorithm calculates a way to tweak the parameters which will reduce the loss function. This tweak is actually a gradient vector, which is only locally accurate. This means that the model will improve if we make a small tweak to the parameters in the direction of the gradient, how big should that adjustment be? Deciding this is the job of the optimiser, and can have a significant effect on how well our model learns. Through a number of tricks such as simulating momentum, and warping the loss landscape to make traversal easier, the ADAM optimiser combines all the most successful optimisation tricks available and is by far the most commonly used. While not guaranteeing the best performance, it can be relied upon to provide consistently good optimisation. Experimenting with other optimisers is an option but not as important as the other considerations in this list.

Validation: Assessing the models performance on the validation dataset regularly during and after training to give an indication of how well the model is learning, and compare between model candidates. Here we need to consider how frequently we want to validate. Start by validating after every epoch. If more insight is required into the learning procedure, increase this to validate after every batch (or n batches), though this will slow down training. * A number of useful tools exist for visualising the validation process. It is critical to do so, to make sure the model is training as expected. **Tensorboard** is highly robust and sufficient for this, and weights and Biases (**WandB**) is a professional subscription service that offers an alternative with a more intuitive user interface and cloud services for synchronised training across multiple devices.

Having done all of this, we are ready to experiment with different model architectures. Model architecture design involves skill, intuition, knowledge, and a not insignificant amount of guesswork. Even the world's top ML researchers could not tell you the optimal number of extra neurons to add to improve model performance without testing it empirically. However, the good news is that these networks are designed to give the best performance they can no matter their design, and you are likely to get reasonably good results within a few attempts. To make this daunting process more systematic, the following section details the key architectural principles that go into the design of a model, and the section after will then lay out a practical guide to follow.

Decision Making Flowchart

Neural Networks

Architectural Principles of Neural Networks

Here we break down some of the key ideas that go into building a more sophisticated neural network than the basic MLP we used in our first example. Understanding these will allow you to see how even the world's most complex and capable neural networks are put together.

- Convolution
 - **What:** A filter with learnable parameters.
 - **Why:** Computer vision has used filters for many decades, designing filters which slide along (convolve with) an image to pick out features such as sharpness, contrast, vertical/horizontal lines etc. The key insight is that if we instead make the parameters in the filter learnable, the algorithm can itself determine what the optimal filters should look like, from the data alone.
- Pooling
 - **What:** A dimensionality reduction technique. Similarly to a convolution, a window is selected on the input, and the corresponding output from this window is the largest activation value found in that window. Then the window is shifted, and the process repeats
 - **Why:** Dimensionality reduction can be desirable for many reasons, but pooling has a number of specific advantages. Firstly, it has no parameters to learn, meaning it adds negligible computational requirements during training. Secondly it has been found to be complementary to convolutions. This is likely because convolutions can be thought of as filters searching for specific features, and the pooling then essentially tells the network if those features are present in the given window.
- Convolutional Blocks
 - **What:** An architectural design shortcut - a combination of convolutional layers, pooling layers and regularisers packaged into a block which is repeated throughout the network.
 - **Why:** Convolutional layers are used to extract local patterns in the input data. By stacking many convolutional layers on top of each other, higher level relationships are able to be recognised. In practice, this means the network is able to perform better (more complex) pattern recognition, which is what machine learning is all about. Using a repeating block structure like the one pictured below has been found to be very effective. Each block contains multiple convolutions and a pooling layer to reduce dimensionality. Each element of the block is optional and customisable, but the principle of repeating blocks remains the same.
- Skip Connections
 - **What:** Passing the output of one layer in the network directly to later layers in the network, 'skipping' over the intermediate layers. This can be done through addition, which requires the dimensions of the layers to be equal, or through appending, which increases the dimension size. Appending is the safer choice, though comes at greater computational cost.

TechNES AI Best Practices Group:

- **Why:** As we make our networks deeper, we are able to extract higher level features. This is extremely powerful, but some new issues begin to emerge. Firstly, the low level information can get lost on the way. Secondly, we can run into the vanishing gradient problem. A neat solution to diminish both of these issues is to use skip connections.
- **Bottlenecks**
 - **What:** A bottleneck refers to the shape of the network (big to small). Often followed by more layers to build the network size back up (small to big). The bottleneck itself is the smallest layer, which can also be called the encoding layer.
 - **Why:** Many networks utilise the idea of a bottleneck, even beyond simple autoencoders (which are nothing more than a bottleneck in structure). Compressing data through a small encoding layer encourages the network to extract the most distinguishing features from the data. This is used in a number of different ways. The encoding itself can be used to represent the data in a unique and low dimensional form. Or the second half of the network can use the information from skip connections and the encoding to infer high level information about the data to solve problems.
- **Recurrence**
 - **What:** A network layer which takes as input some data and a 'state' vector, and produces a new state vector alongside its other output.
 - **Why:** The state vector represents some understanding about the state at a given time. The idea then is for the layer to take in some data and update this understanding, so that the state is different for the next time step. Without recurrence, networks have no notion of time or way to relate the data coming in now with what came before. This is necessary for sequential tasks like video or text recognition, and not necessary for static tasks such as image recognition, hence the discrepancy.
- **Attention**
 - **What:** Weighting each element of a sequence of data, according to how important (how much attention should be paid to) it, to solve a given task. Typically, the model will be outputting another sequence, and as such each element of the output will require a different set of attention weights. The full theory behind attention is beyond the scope of this guide. It is listed here to give a brief intuition behind transformer models, which are growing in popularity and based on the principle of attention.
 - **Why:** Attention provides a way for a network to take in sequential data all at once, learning which input elements each output should pay attention to. This offers numerous benefits over recurrence, such as the ability to be processed in parallel (recurrent networks are inherently sequential so cannot be parallelised), and removing recency bias. Recency bias being the tendency of RNNs to pay more attention to the most recent sequence elements, rather than the most relevant.

Neural Network Design and Experimentation Process

1. Establish a strong baseline. The first thing to do once your data pipeline and evaluation metrics are set up, is to try out the simplest neural network design relevant to your problem. Usually this will be an MLP. The data will determine the input and output size, so for this make a simple MLP with one hidden layer. There is no magic formula to tell you how large to make this hidden layer, the key is to experiment. Start with $(\text{input size} + \text{output size})/2$ if you are unsure. Train and evaluate. This gives you a benchmark and some idea of how complex the task will be. If the MLP performs very well, you may wish to stop there. If not, move on to step 2.
2. Design a neural network specifically for your problem. How? Google it! More specifically, scour the internet for a research paper, article or competition submission that publishes a machine learning model for a problem similar to yours. All kinds of similarity are useful, working on the same data type

TechNES AI Best Practices Group:

(eg time series, video etc), or the same problem (eg anomaly detection, object recognition), but ideally both. For this to be successful, it is likely that the authors have already done a great deal of the work for us in choosing approximately the right kind of network architecture. Take this as a starting point.

- a. The closer their problem is to yours, the less we need to experiment with other architectures
 - b. For research papers in particular, models may be more complicated than necessary as authors are usually proposing a novel method. We can experiment with removing the more complicated features, if they don't affect performance.
 - c. If nothing relevant arises, use the baseline model we established earlier as the starting point.
3. Iteration and experimentation. Given a baseline (either from step 1 or 2) and working data pipeline, it may be surprisingly straightforward to test different ideas, so long as enough computational resources are available. Bear your initial goals in mind, don't be afraid to stop iterating when these are met even if it may be possible to squeeze slightly more performance out with further experimentation. Very small improvements to measured performance on test data may not actually translate to a significantly better model in the real world. APIs such as TensorFlow and PyTorch make adjusting model architectures as simple as stacking preset functions on top of one and other. These all come in built with these APIs, and are listed below:

- a. Regularisers: These may improve generalisation performance without changing the overall architecture.

- i. Dropout

- ii. Batch normalisation.

- b. Dimensionality Manipulations: Scale up or down dimensionality. Often comes at a cost of information loss in favour of computational efficiency.

- i. Max Pooling

- c. **Recurrent modules: for sequential data only.**

- i. LSTM

- ii. GRU

- d. **Other**

- i. Convolutions

- ii. Skip Connections

Taking the ideas of others as inspiration, try out some of these ideas and observe their effects, both individually and in combination. See the architectural principles section for a breakdown of what each of these do. Remember, simplicity is key. A useful method to avoid wasted time is to train and test the model after a given small change, one change at a time. If performance doesn't improve, don't waste any more time on changes of that kind. Whereas if you make many changes and then evaluate, you can't be sure which changes are having a positive effect.

Worked Examples

Pitfalls and Common Problems

Resources

Appendix